

Introduction to Deep Learning

Mnacho Echenim

Grenoble INP-Ensimag

2022-2023



Notes

About this class

- Deep learning
 - ▶ Very dynamic field of research
 - ★ Increased computational power
 - ★ Increased data availability
 - ▶ Different kinds of neural networks, depending on applications
 - ★ Convolutional networks for image processing
 - ★ Recursive neural networks, large language models for natural language processing
 - ★ Generative adversarial network for physical simulations and image processing
 - ▶ But not many impressive results obtained by using a neural network as a black box, with no knowledge on how they work or on the data
 - ★ Other ML techniques can outperform neural networks, depending on the problem at hand
- Goal of this class
 - ▶ Present the fundamentals of deep learning by implementing an efficient Multilayer Perceptron (MLP) from scratch
 - ▶ Use this MLP to solve a regression problem (**supervised learning**)
 - ▶ **Nb:** you may feel lost at the beginning of the hands-on labs but everything will fall into place with time



Notes

Main outline

- Origins of neural networks and deep learning
- Main properties of neural networks
- Learning techniques
 - ▶ Basis (gradient descent, overfitting, ...)
 - ▶ Improvements (gradient optimization, regularization, ...)
- Application to a pricing problem
 - ▶ Question: can a neural network be used as a proxy for a costly function?

Notes



Timetable

- 1 Introduction, multilayer perceptrons, forward propagation
 - ▶ Hands-on: forward propagation
- 2 Gradient descents, backpropagation
 - ▶ Hands-on: backpropagation for stochastic gradient descent
- 3 Mini-batch gradient descent, gradient descent optimization
 - ▶ Hands-on: mini-batch gradient descent
- 4 Hands-on: mini-batch, Momentum
- 5 Regularization
 - ▶ Hands-on: L2 regularization
- 6 Feature preprocessing, batch normalization
 - ▶ Pricing project
- 7 Energy considerations
 - ▶ Pricing project
- 8 Pricing project

Notes



Some resources

- Goodfellow, Bengio, Courville. Deep Learning:
<https://www.deeplearningbook.org/>
- Aggarwal. Neural Networks and Deep Learning
- Nielsen. Neural Networks and Deep Learning:
<http://neuralnetworksanddeeplearning.com>
- Lectures by Mallat at Collège de France:
<https://www.college-de-france.fr/site/stephane-mallat/>
- Lectures by Le Cun at Collège de France:
<https://www.college-de-france.fr/site/yann-lecun/>
- Stanford School of Engineering channel:
<https://www.youtube.com/user/stanfordeng>

Notes

Some landmarks

- 1943 W. McCulloch, W. Pitts: *A Logical Calculus of the Ideas Immanent in Nervous Activity*
- Formal model of a neuron
 - Weighted inputs, binary outputs
- 1948 N. Wiener: *Cybernetics: Or Control and Communication in the Animal and the Machine*
- “Self-regulating mechanisms”
 - Notion of feedback
- 1950 A. Turing: The Turing Test
- Proposal: build a program simulating a child’s mind, then educate it
- 1951 M. Minsky and D. Edmonds: Stochastic Neural Analog Reinforcement Calculator (SNARC)
- First neural network machine

Notes

A classification problem

1958. F. Rosenblatt: *The Perceptron: A Probabilistic Model for Information Storage and Organization in the Brain*

- Input: $\{(\chi_{(i)}, \rho_{(i)}) \mid i = 1, \dots, N\}$,
where $\chi_{(i)} \in \mathbb{R}^d$ and $\rho_{(i)} \in \{+1, -1\}$
- Goal: construct a classifier that correctly labels the training set
- Perceptron: outputs $\text{sign}(\omega^T \chi + \beta)$
- Problem: find ω, β such that
 $\forall i = 1, \dots, N,$
 $\text{sign}(\omega^T \chi_{(i)} + \beta) = \rho_{(i)}$
- Reformulation: find ω, β such that
 $\forall i = 1, \dots, N, \rho_{(i)} \cdot (\omega^T \chi_{(i)} + \beta) > 0$

Notes

Parameter update algorithm

Notations

We define:

- $\theta \stackrel{\text{def}}{=} (\beta, \omega_1, \dots, \omega_d)^T$ and $\bar{\chi} \stackrel{\text{def}}{=} (1, \chi_1, \dots, \chi_d)^T$

Goal: find θ such that $\forall i = 1, \dots, N, \rho_{(i)} \cdot (\theta^T \bar{\chi}_{(i)}) > 0$

Notes

Input: $\{(\chi_{(i)}, \rho_{(i)}) \mid i = 1, \dots, N\}$

```
1  $\theta \leftarrow (0, \dots, 0)^T$ ;  
2 while  $\exists j \in \{1, \dots, N\}$  such that  $\rho_{(j)} \cdot \text{sign}(\theta^T \bar{\chi}_{(j)}) < 0$  do  
3   |  $\theta \leftarrow \theta + \rho_{(j)} \bar{\chi}_{(j)}$   
4 end  
5 return  $\theta$ 
```

Convergence

Theorem

If there exists a separating hyperplane for the samples, then the algorithm terminates and the resulting perceptron correctly classifies all samples

- Separating hyperplane: $\exists \bar{\omega}^*$ such that for all $i = 1, \dots, N$, $\rho_{(i)} \cdot (\bar{\omega}^{*T} \bar{\chi}_{(i)}) > 0$
 - ▶ Note that $\bar{\omega}^* \neq \mathbf{0}$
- Proof principle: determine an upper-bound on the number k of iterations in the **while** loop
- We let:
 - ▶ $\delta \stackrel{\text{def}}{=} \min_i [\rho_{(i)} \cdot (\bar{\omega}^{*T} \bar{\chi}_{(i)})]$; note that $\delta > 0$
 - ▶ $R \stackrel{\text{def}}{=} \max_i \|\bar{\chi}_{(i)}\|$; note that $R > 0$
- We have (See, e.g., <http://www.cs.columbia.edu/~mccollins/courses/6998-2012/notes/perc.converge.pdf>):

$$k \leq \frac{R^2 \|\bar{\omega}^*\|^2}{\delta^2}$$

Notes

What if the samples admit no separating hyperplane?

Minsky & Papert (1969): *Perceptrons: an introduction to computational geometry*

- The XOR function does not admit a separating hyperplane
- The perceptron is **incapable** of learning this function

Notes

Summary on the perceptron

- Simple learning scheme
- Convergence result when there exists a separating hyperplane
 - ▶ It is possible to evaluate the convergence speed
- Minsky & Papert: there are simple functions that cannot be learned
 - ▶ This negative result slowed research on neural networks for a long time
- Question: what happens if we generalize the perceptron by
 - ▶ Interconnecting several neurons together
 - ▶ Allowing activation functions other than the sign function?

Notes

Definitions

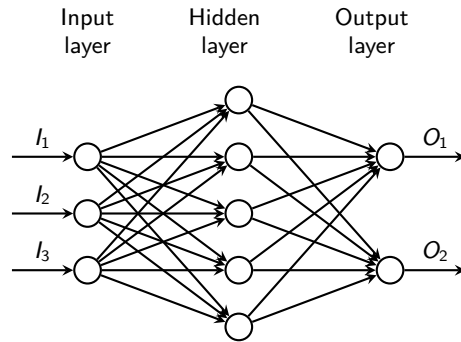
Definition

A multilayer perceptron (MLP, or feedforward neural network) is a set of interconnected neurons that forms a Directed Acyclic Graph (DAG)

- The connections between neurons are **weighted**
- Neurons with no incoming connection from another neuron are **input neurons**. The set of input neurons forms the **input layer**
- Neurons with no outgoing connection to another neuron are **output neurons**. The set of output neurons forms the **output layer**
- All other neurons are **hidden neurons**. Hidden neurons are organized in **hidden layers**
- Computations are performed in the hidden and output layers.

Notes

Fully connected multilayer perceptron (3 layers)



Definition

A fully connected MLP is an MLP such that all non-input neurons are connected to all the neurons of the previous layer

In what follows, all the MLPs we consider will be fully connected

Notes

Notations and assumptions

Given a fully connected MLP with L layers:

- The neurons of layer i are denoted by $\nu_1^i, \dots, \nu_{n_i}^i$ (layer 0 denotes the input layer)
- The weight connecting ν_j^{i-1} to ν_k^i is denoted by $\omega_{j,k}^i$
- The bias of ν_k^i is denoted by β_k^i
- The sample transmitted to the input layer is denoted by $\alpha^0 = (\alpha_1^0, \dots, \alpha_{n_0}^0)^T$
- The **net input** of ν_k^i is $\zeta_k^i \stackrel{\text{def}}{=} \left(\sum_{j=1}^{n_{i-1}} \omega_{j,k}^i \alpha_j^{i-1} \right) + \beta_k^i$
- The **activation** of ν_k^i is denoted by α_k^i
- The activation function of ν_k^i is Φ , so that $\alpha_k^i = \Phi(\zeta_k^i)$
 - ▶ Unless stated otherwise (e.g., **Dropout**), we assume it is the same for all neurons

Assumption

In what follows, unless specified otherwise, we will assume that the MLPs we consider have a single output node

Notes

Matrix representation

- The weights at layer i are stored in a matrix $\Omega^i \in \mathbb{R}^{n_{i-1} \times n_i}$, where $(\Omega^i)_{j,k} = \omega_{j,k}^i$
- The biases at layer i are stored in a vector $\beta^i \stackrel{\text{def}}{=} (\beta_1^i, \dots, \beta_{n_i}^i)^T$
- The net inputs at layer i are stored in a vector $\zeta^i \stackrel{\text{def}}{=} (\zeta_1^i, \dots, \zeta_{n_i}^i)^T$
- The activation function is extended to vectors: if $u = (u_1, \dots, u_n)^T$, then $\Phi(u) = (\Phi(u_1), \dots, \Phi(u_n))^T$
- The activations of layer i are stored in a vector $\alpha^i \stackrel{\text{def}}{=} (\alpha_1^i, \dots, \alpha_{n_i}^i)^T$

Notes

Forward propagation algorithms

Input: A network with L layers

Input: $\alpha^0 \stackrel{\text{def}}{=} (\alpha_1^0, \dots, \alpha_{n_0}^0)^T$

```
1 for i ← 1 to L do
2   for k ← 1 to ni do
3      $\zeta_k^i \leftarrow \left( \sum_{j=1}^{n_{i-1}} \omega_{j,k}^i \alpha_j^{i-1} \right) + \beta_k^i$ ;
4      $\alpha_k^i \leftarrow \Phi(\zeta_k^i)$ ;
5   end
6 end
```

Algorithm 2: Forward propagation, basic version

Notes

Identity activation function

What happens if the activation function is the identity?

- $\alpha^i = \zeta^i = [\Omega^i]^T \alpha^{i-1} + \beta^i$

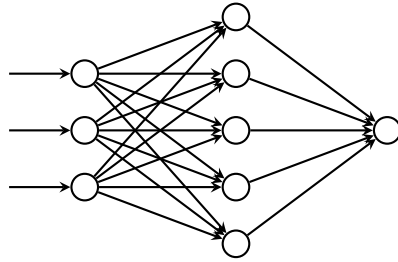
Notes

Linear activation function

What happens if the activation function is linear: $\Phi : z \mapsto \mu z + \lambda$?

Notes

A specific MLP



- Input layer of size d
- Single hidden layer of size K
- Single output neuron
- Activation function Φ only on hidden layer (identity on output)
- No bias on output layer

Notes

A class of functions

- We define $\Lambda \stackrel{\text{def}}{=} \{ \langle K, \Phi, \Omega, \omega', \beta \rangle \mid K \in \mathbb{N}, \Phi : \mathbb{R} \rightarrow \mathbb{R}, \Omega \in \mathbb{R}^{d \times K}, \beta, \omega' \in \mathbb{R}^K \}$
- For $\lambda \in \Lambda$, we define

$$\begin{aligned} f_\lambda : \mathbb{R}^d &\rightarrow \mathbb{R} \\ x &\mapsto [\omega']^T \cdot \Phi(\Omega^T x + \beta) = \sum_{k=1}^K \omega'_k \Phi([\Omega_k]^T x + \beta_k) \end{aligned}$$

- We consider the set of functions $\mathcal{F}(\Lambda) \stackrel{\text{def}}{=} \{ f_\lambda \mid \lambda \in \Lambda \}$

Question

What functions can be approximated by an element from $\mathcal{F}(\Lambda)$?

Notes

The Universal Approximation Theorem

Theorem (see, e.g., Allan Pinkus, 1999)

Let $\mathcal{C}(\mathbb{R}^d)$ denote the set of continuous functions on \mathbb{R}^d and consider the set $\mathcal{M}(\Phi) \stackrel{\text{def}}{=} \text{span}(\{x \mapsto \Phi(\omega^T x + \beta) \mid \omega \in \mathbb{R}^d, \beta \in \mathbb{R}\})$. This set is dense in $\mathcal{C}(\mathbb{R}^d)$, in the topology of uniform convergence on compact sets if and only if Φ is nonpolynomial

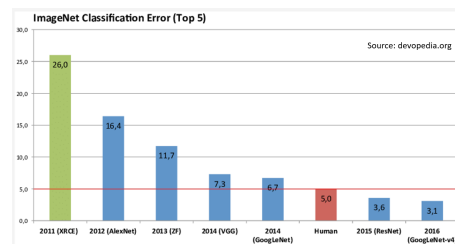
In other words: if $f \in \mathcal{C}(\mathbb{R}^d)$ and $\mathcal{D} \subseteq \mathbb{R}^d$ is a compact set, then for all $\varepsilon > 0$, there exists $\tilde{f} \in \mathcal{M}(\Phi)$ such that for all $x \in \mathcal{D}$, $|f(x) - \tilde{f}(x)| \leq \varepsilon$

- For a given $\varepsilon > 0$, there exists $K_\varepsilon \in \mathbb{N}$ such that an MLP with a single hidden layer of size K_ε can approximate f with a precision of ε
- What is wrong with polynomials?

Notes

Issues with the Universal Approximation Theorem

- The size of the hidden layer can be extremely large
 - ▶ In practice, this result does not seem very useful
- What happens if we increase the number of hidden layers and bound their sizes?
 - ▶ For a long time, not so clear there was anything to gain
 - ▶ Empirically yes: impressive results using deep networks



- ▶ More recent theoretical results (Eldan, Shamir. The power of Depth for Feedforward Neural Networks, 2016)
- ▶ For the time being, it is still not possible to formally explain why things work so well

Notes
